

CHRON — REALITY MAP

What CHRON actually is on this machine, right now. Every line below is answerable to an artifact that was opened, hashed, counted, or executed during this audit. Nothing is inferred from a declaration.

REFERENCE	AUTHORITY	MODE	MUTATIONS
CHRON::REALITY_MAPPER::v1	ARIF (F13)	OBSERVE_ONLY / AUDIT	0 — read-only
RENDER	SEAT	CLASS	VERDICT MODEL
2026-09-18 12:11 MYT (04:11Z)	HERMES (i-ARIF) on KVM8	Lane B RECEIPT — not a SEAL	UNKNOWN / SABAR / HOLD / PARTIAL / SEAL

The constraint this document is written under. DECLARATION ≠ MECHANISM. A file named `chron.py` is not a running clock. A job marked `enabled` is not a job that fired. An empty namespace is not a quiet subsystem. Where the evidence does not exist, this map says `UNBUILT` or `UNKNOWN` — never `PASS`. The federation's own spine gate already sets that precedent: a boundary with no artifact to test it reports `UNBUILT`, "honest, and NOT a pass."

Primary question

"Apakah task, loop, capability, dan reality-flow CHRON yang sedang hidup sekarang?"

Answer in one sentence, from evidence only: **a date store (5 events), a sealed gate script (6 PASS / 3 UNBUILT), one live card lane that has produced two artifacts and delivered none, a growing receipt pile in which 100% of rows lack a trace_id, and a learning loop that inhales twice a day and has closed zero policies.** Everything else that carries the name CHRON today is either a document or an empty directory.

PHASE 1 — REALITY INVENTORY

Classes are fixed by the mission. Confidence: HIGH = opened/counted/executed this audit; MED = surface confirmed, contents not exhaustively read; LOW = name only.

1.1 EVENT SOURCES

ID	PATH / SURFACE	LAST ACTIVITY	DECLARED	OBSERVED	CONF
E1	/root/.hermes/cron/jobs.json	2026-09-18T12:06:12+08	job registry	39 jobs · 13 enabled · updated today	HIGH
E2	/root/.hermes/cron/deliveries.db	2026-09-18T12:06	delivery ledger	rows delivered today (12:06, 09:59, 07:02)	HIGH
E3	systemctl list-timers --all	2026-09-18T12:11	substrate health	30+ timers live; none named chron	HIGH
E4	root crontab	2026-09-18 (hourly jobs firing)	infra loops	~45 lines active; no CHRON entry	HIGH
E5	/etc/cron.d/aaa-rsi-loop	2026-09-18T12:07:02+08	RSI exhale loop	present + firing (see L1-L4)	HIGH
E6	/root/AAA/scripts/chron_events.json	2026-09-18T10:25	canonical date store	5 events, 5/5 carry source+confidence+date+audience	HIGH
E7	forge_work/alpha-zen/cycles.jsonl	2026-09-18T10:40	render receipts	8 cycle records; 6 render_status OK; last 1 rejection_tracking=NOT_MEASURED	HIGH
E8	~/.local/share/arifos/event_bus.jsonl	2026-09-17T22:15:02Z	event spine	last record sealed_events_sweep → DEFECTS_FOUND exit 1	MED
E9	~/.local/share/arifos/ariffow_receipts.jsonl	2026-09-18T03:51:44Z	execution receipts	35,743 rows · 35,743/35,743 trace_id NULL	HIGH

1.2 MEMORY SOURCES

ID	PATH / SURFACE	LAST ACTIVITY	DECLARED	OBSERVED	CONF
M1	/root/AAA/claim_ledger/claims.db	2026-09-18T11:51	claim ledger	21 claims · 11 verifications · 6 artifacts · 38 chain rows · append-only triggers present. Separate organ — not CHRON-owned.	HIGH
M2	A-FORGE/duties/logs/world-model-lite.jsonl	2026-08-27T02:20:26Z	world model	792 rows parsed · last write 22 days ago · writer is its only reader	HIGH
M3	Qdrant :6333	live at probe	vector memory	20 collections (arifos_memory, vault_canon, atlas333_eureka...) — presence only	MED
M4	/root/arif05/VAULT999/chronus/	dir mtime 2026-09-09	CHRON vault namespace	EMPTY — 0 entries, 0 bytes	HIGH
M5	forge_work/alpha-zen/cards/	2026-09-18T10:32	episodic card store	2 records (morning/night), 1 day deep	HIGH
M6		(directory present)	Reality Graph	existence confirmed; contents not enumerated this pass	UNKNOWN

1.3 PREDICTION SOURCES

ID	PATH / SURFACE	LAST ACTIVITY	DECLARED	OBSERVED	CONF
P1	~/.local/share/arifos/state/closure-ledger.jsonl	2026-09-17T23:45	prediction store	4 records carrying {claim, predicted, verify_cmd, verify_by, observed, verdict} — the only live prediction object found	HIGH
P2	world-model-lite.jsonl	2026-08-27T02:20:26Z	prediction stream	792 predictions, 792 distinct ids, 5 distinct actions , allow_mutate=False on 734. No outcome / verify_at field exists in the row schema.	HIGH
P3	calibration records	—	Brier / log-loss / ECE	NOT FOUND. Only a proposal: AAA/scripts/proposals/apex-zen-calibration-2026-09-13.md	HIGH

1.4 ATTENTION SOURCES

ID	PATH / SURFACE	LAST ACTIVITY	DECLARED	OBSERVED	CONF
A1	/root/scripts/attention_governor.py (21 KB)	2026-09-16T11:00	attention governor	present; scans carry / cron / repos / organs / pointers / seal	HIGH
A2	state/attention-ledger.jsonl	2026-09-17T23:45	attention ledger	9 rows only. 09-16: 32 loops. 09-17: 30 loops, composition changed (see D6)	HIGH
A3	state/attention-2026-09-17.json	2026-09-17T23:45	daily snapshot	16.8 KB, status: COMPLETE	HIGH
A4	Hermes job attention-closure (23:45)	2026-09-17T23:46 ok	daily closure pass	enabled · -no-agent · empty stdout delivers nothing	HIGH
A5	ALPHA-ZEN cards + docforge editions	2026-09-18T10:32 / 09:57	human-facing cards	2 card artifacts today; EDITION-001.pdf + EDITION-002 built with light-theme + ink gates	HIGH

1.5 LEARNING SOURCES

ID	PATH / SURFACE	LAST ACTIVITY	DECLARED	OBSERVED	CONF
L1	AAA/rsi/state/atoms.jsonl	2026-09-18T12:07	RSI atoms	346 atoms · 318 verified · 28 rejected	HIGH
L2	AAA/rsi/state/loop-ledger.jsonl	2026-09-18T12:07	loop receipts	14 cycles · 5 EXHALED · 8 INHALE_ONLY · 1 note. Latest 4 cycles all INHALE_ONLY	HIGH
L3	AAA/rsi/state/consequence.jsonl	2026-09-18T12:07	consequence ledger	33 rows · 7 PENDING (consequence not yet observable)	HIGH
L4	AAA/rsi/state/baselines.json	2026-09-18T12:07	capability baselines	7 baselines, all status PENDING , 7-day windows	HIGH
L5	AAA/governance/UNRATIFIED-LESSONS-LEDGER.jsonl	(12 rows)	lesson ledger	10 UNRATIFIED_LESSON + 2 PENDING_V12_FOLD → 0 ratified	HIGH
L6	AAA/scars/	last new scar 2026-09-15	scar store	13 scars + 4 auto-candidates (latest 2026-09-15)	HIGH
L7	POLICY CHANGES	—	lesson → policy promotion	0 OBSERVED. No policy-candidate object exists; nothing can be promoted from one episode	HIGH

Phase 1 finding — the inventory's shape

CHRON's **evidence** is real and countable on every surface above. What is not present is CHRON **as a running thing**: there is no service, no timer, no cron line, and no job whose name or target is CHRON. The nearest live neighbours each own a different organ — the RSI loop belongs to capability evolution, the attention governor belongs to the attention lane, docforge-edition-daily belongs to the briefing lane. CHRON's own code exists and runs on demand; nothing schedules it.

PHASE 2 — SPINE EXTRACTION

The spine below is the one CHRON itself declares and tests: WORLD → OBSERVE → ENCODE → SELECT → DELIVER → RESPONSE → OUTCOME → COMPARE → EXPERIENCE → MEMORY → LEARNING → SHADOW → VALIDATION → PROMOTION → NEW VERSION. It was **executed** this audit — python3 /root/AAA/scripts/chron_spine_gate.py — not read. Raw verdict: **6 PASS · 0 FAIL · 3 UNBUILT**, self-test PASS (every check can reject).

2.1 Edges with an artifact behind them

EDGE	SOURCE → TARGET	DECLARED	OBSERVED	PROOF	LAST SEEN	STATUS
E-01	WORLD → OBSERVE	events carry provenance	5/5 events carry source, confidence, date, audience	chron_events.json + gate check 1	2026-09-18T10:25	PASS
E-02	OBSERVE → CLAIM	every cell classed	9 rows, every cell tagged (OBS/DER/INT/CONTESTED)	gate check 2 + cards/2026-09-18-*.json	2026-09-18T10:32	PASS
E-03		truth ≠ urgency		gate check 3	2026-09-18T11:02	PASS

	CLAIM → ATTENTION		truth and urgency exposed separately; provably diverge on a collision pair			
E-04	ATTENTION → DELIVERY	newest row resolves	newest row resolves; 4 older rows stale, all predating the content-addressing fix	gate check 4	2026-09-18T11:02	PASS
E-05	DELIVERY → HUMAN	one reader contract	12 patterns (10 word-bounded prose, 2 literal); inner-state rule present	gate check 5	2026-09-18T11:02	PASS
E-06	LEARNING → PRODUCTION	learner behind a ladder	engine defaults to SHADOW and cannot send; a separate production renderer exists — the ladder holds	gate check 8	2026-09-18T11:02	PASS
E-07	EXPERIENCE → MEMORY	episodes persist	collection has no payload schema	gate check 6	–	UNBUILT
E-08	MEMORY → POLICY	lesson promotion	no policy-candidate object exists	gate check 7	–	UNBUILT
E-09	OUTCOME → LEARNING	verified prediction retained	no prediction object exists in that lane, so no failed prediction can be retained	gate check 9	–	UNBUILT

2.2 Edges the declaration implies but no artifact carries

EDGE	DECLARED	OBSERVED	PROOF	STATUS
ENCODE (format contract)	canonical encode step	alpha_zen_card.schema.json exists and gates a card; no CHRON-owned encode contract	schema file present, not referenced by any CHRON trigger	PARTIAL
RESPONSE (did the human reply?)	human response is an event	delivery receipts exist; no response-capture edge from the reader back into CHRON	deliveries.db shows DELIVERED, never REPLIED	UNKNOWN
COMPARE (predicted vs observed)	outcome comparison	2 of 4 closure-ledger records reached a verdict; the rest open or contradicted	closure-ledger.jsonl	PARTIAL
SHADOW / VALIDATION / PROMOTION	maturity chain before promotion	RSI verifier marks itself independence_class: NOMINAL · independent: false · provisional: true	rsl/state/receipts/loop-2026-09-18T120702+0800.json	PARTIAL
NEW VERSION	promotion produces a new version	0 policies promoted , so no new version has been produced	UNRATIFIED-LESSONS-LEDGER (12 rows, 0 ratified)	FAIL

PHASE 3 — DECLARATION vs REALITY (REALITY_DRIFT_LEDGER)

#	CAPABILITY	DECLARED_MECHANISM	OBSERVED_MECHANISM	CONTRADICTION	SEV
D1	ALPHA-ZEN three-pulse lane	3 jobs enabled, delivering cards to the two-person room -1003815535761 at 07:15 / 14:00 / 21:15	All 3 created 2026-09-18T10:02 ; last_run_at = null on all 3; repeat.completed = 0	Jobs were born 2h47m after their own 07:15 slot. A card pipeline that pass all its gates has zero executions and zero delivery receipts . SCHEDULED ≠ RAN.	HIGH
D2	Daily sovereign brief	docforge-edition-daily · 0 6 * * * → PDF to Arif's DM	Created 09:34 today; first run 09:59 (EDITION-001.pdf, real PDF, gates passed) — nearly 4 h after the declared window	The 06:00 brief has never run at 06:00. For 2026-09-18 no 06:00 artifact exists at all; what exists is a 09:59 artifact under a 06:00 label.	MED
D3	World model / prediction substrate	a maintained prediction layer	792 predictions, 792 distinct ids, only 5 distinct actions ; last write 2026-08-27T02:20Z; row schema has no outcome field ; the writer is the only reader	Artifacts look healthy (valid JSONL, growing to 600 KB) while the producer silently ceased 22 days ago . Exactly the failure shape the federation already named: a subsystem whose artifacts look fine while it has stopped.	HIGH
D4	VAULT999 CHRON namespace	VAULT999/chronus/ exists as a namespace	directory present, 0 entries	A named namespace with zero records reads as 'nothing to seal' when it actually means 'nothing was ever sealed'. Empty ≠ quiet.	MED
D5	Execution receipts as causal ledger	receipts bind action to objective	35,743 rows; 35,743 rows with trace_id = NULL (100%); last row verdict UNKNOWN · signal FQ_SIGNAL_DRIFT	A receipt without a trace_id is an event-pile entry, not a causal-ledger entry. The doctrine that named this defect is being violated by the running system, measurably, at 100%.	HIGH
D6	Attention closure (silent jobs)	CRON_SILENT is tracked and surfaced	2026-09-16 ledger: 32 loops — CRON_DISABLED 15, CRON_SILENT 6. 2026-09-17: 30 loops — CRON_DISABLED 22, CRON_SILENT absent	The silent-job category did not get fixed; it disappeared while the disabled count rose 15 → 22. A reclassification, not a resolution — and the closure ledger's own record (REALITY_DISAGREED) proves reality refused the prediction.	MED
D7	KVM4 cron reachability	cron API not exposed on KVM4		The recorded finding said the endpoint 'returns empty'. It does not — it refuses	MED

			GET /cron/jobs → HTTP error proxy_attribution_required (not empty, not 404)	with a proxy-attribution error. 'No data' was read as 'all clear'. /health answers {ok:true,status:live} on the same port.	
D8	Audit artifact currency	blockers B1-B4 describe live code	alpha_zen_engine.py was repaired at 11:18 today — the year-wrong budget date and the frozen od1_days literal are gone from the code; AUDIT-ALPHA-ZEN-2026-09-18.json records no supersession	The blockers are fixed in the code and still stand in the audit. A reader picking up that artifact tomorrow inherits four defects that no longer exist — Stale Claim Propagation, in the artifact whose job was to prevent exactly this.	MED

Phase 3 note on the one contradiction that went the other way

Not every discrepancy is decay. chron.py and alpha_zen_engine.py now **read one canonical store** (chron_events.json) and compute every day-count at render; the engine returns **no countdown at all** rather than inventing one when the store is unreadable. That is the correct repair of the exact defect class this phase hunts — two clocks drifting silently because both numbers look plausible. Recorded here so the ledger is not one-sided.

PHASE 4 — TASK EXTRACTION (only tasks with evidence)

No task below is proposed. Each is either a registered trigger observed in the scheduler, or a coded pipeline observed on disk and clearly wired to one.

TASK_ID	PURPOSE	TRIGGER	OUTPUT	SCHEDULE	LAST_RUN	LAST_SUCCESS	DEPENDENCIES	RECEIPT_PATH	STATUS
CH-T01	Morning dual-lens card → two-person room	Hermes cron 8313453a73a7	1 message + ALPHA-ZEN-MORNING.png	15 7 * * *	null	—	skills/apex-cognitive-reflex/alpha-zen-lanes · chron.py · alpha_zen_card.py	cards/2026-09-18-morning.json	ARM
CH-T02	Afternoon body check → same room	Hermes cron 3e40c50f0ca2	1 message	0 14 * * *	null	—	WELL organ tools · chron.py	cycles.jsonl (manual runs only)	ARM
CH-T03	Night card → same room	Hermes cron 161c0d5e0d0c	1 message + ALPHA-ZEN-NIGHT.png	15 21 * * *	null	—	GEOX · WEALTH · chron.py · alpha_zen_card.py	cards/2026-09-18-night.json	ARM
CH-T04	Sovereign daily brief PDF → Arif DM	Hermes cron a77681618199	EDITION-00N.pdf + ledger + receipt	0 6 * * *	2026-09-18T09:59:06	ok (3h59m late)	weasyprint 69.0 · base-a4 template · brief-rules-context.py	forge_work/docforge-ledger.jsonl	ACT
CH-T05	Daily attention closure + quiet scan	Hermes cron attention-closure	ledger row + snapshot (silence if clean)	45 23 * * *	2026-09-17T23:46	ok	attention_governor.py · closure_probe.py	state/attention-ledger.jsonl	ACT
CH-T06	Capability evolution (inhale→exhale)	/etc/cron.d/aaa-rsi-loop	atoms, loop receipt, baselines	every 6 h @ :07	2026-09-18T12:07	INHALE_ONLY (no exhale)	promote.py (FORBIDDEN_PATHS self-test) · rsi verifier	rsi/state/receipts/loop-*.json	ACT
CH-T07	FI mesh health sentinel	/etc/cron.d/aaa-mesh-health	mesh-health.json + holds	17 3 * * *	(daily)	—	mesh-health-probe.sh	forge_work/mesh-health.log	ACT
CH-T08	Prediction stream for federation state	A-FORGE duties loop	world-model-lite.jsonl rows	declared continuous	2026-08-27T02:20Z	stale 22 d	world_model_lite.py — no consumer	duties/logs/world-model-lite.jsonl	BRO
CH-T09	Gate + self-test over the spine	manual / audit only	6 PASS · 0 FAIL · 3 UNBUILT	no trigger	2026-09-18T12:10 (by this audit)	PASS	chron_events.json · cards/ · cycles.jsonl	stdout only — no receipt path	DOR
CH-T10	Render the CHRON clock	manual / called from job prompts	ranked temporal list to stdout	no trigger	invoked inside CH-T01..03 prompts	PASS	chron_events.json	none	DOR

Phase 4 finding — armed is not fired

Three tasks (CH-T01..03) carry a full contract, a named audience, a gate that can refuse, and a rendered artifact already sitting on disk from manual runs. None has ever executed. The difference between that and a working lane is not code — it is **one receipt**. The next scheduled opportunity is **14:00 MYT today** (CH-T02).

PHASE 5 — LOOP CLOSURE AUDIT (9 loops)

#	LOOP	EXISTS	PROOF	LAST EVIDENCE	STATUS
L1	Observation → Witness	Yes	5/5 events carry provenance; spine gate check 1	2026-09-18T10:25	PASS
L2	Witness → Claim	Yes		2026-09-18T11:51	PASS

21 claims with source_ref + locator + sha256 binding; append-only triggers; spine check 2

L3	Claim → Prediction	Partly	Only closure-ledger carries a prediction object — 4 records total ; none derived from the claim ledger automatically	2026-09-17T23:45	PARTIAL
L4	Prediction → Verification	Yes, small	verify_cmd + verify_by present on every closure row; appointment honoured at 23:45	2026-09-17T23:45	PASS
L5	Verification → Outcome	Yes	Observed values recorded: 2 × PREDICTION_MATCHED, 1 × REALITY_DISAGREED, 1 OPEN due 2026-09-21	2026-09-17T23:45	PARTIAL
L6	Outcome → Calibration	No	No Brier / log-loss / ECE artifact anywhere; scoring never computed; only a 2026-09-13 proposal	—	FAIL
L7	Error → Lesson / Scar	Yes	12 unratified lessons + 13 scars + 4 auto-candidates; 28 RSI atoms rejected rather than absorbed	2026-09-18T12:07	PASS
L8	Lesson → Policy	No	12 lessons, 0 ratified , 0 policy-candidate objects; RSI organs outcome reads applied: false, reason: deferred	—	FAIL
L9	Policy → Attention	No	Attention is generated from <i>gaps</i> (stale pointers, silent cron, dirty repos), not from policy; no policy→attention edge observed	2026-09-17T23:45	FAIL

Phase 5 finding — the loop closes at verification and dies there

Six of nine loops have real artifacts. Three do not, and they are the three that would make CHRON *learn*: calibration, promotion, and attention feedback. The system today can **observe, claim, predict, and verify**. It cannot **score itself, promote a lesson, or let a promoted policy change what it looks at next**. That is the learning-closure problem, and it is measured here rather than asserted: one failed prediction (REALITY_DISAGREED, silent-job count = 6) exists on disk and has produced no calibration entry, no lesson row, and no policy.

PHASE 6 — CHRON MATURITY SCORE (counts only, no subjective scoring)

6.1 Observed totals — each figure bound to the artifact that produced it

METRIC	COUNT	SOURCE ARTIFACT	STATE OF THAT ARTIFACT
TOTAL_EVENTS_OBSERVED	14 named events (5 CHRON store + 8 render cycles + 1 event-bus record)	chron_events.jsonl · cycles.jsonl · event_bus.jsonl	last event-bus record ends DEFECTS_FOUND exit 1
TOTAL_RECEIPTS_OBSERVED	35,743 (ariffow) + 38 (claim chain) + 9 (attention)	ariffow_receipts.jsonl · claims.db::ledger_chain · attention-ledger.jsonl	100% of the 35,743 carry trace_id = NULL
TOTAL_PREDICTIONS_OBSERVED	796 (792 world-model + 4 closure-ledger)	world-model-lite.jsonl · closure-ledger.jsonl	792 are dead & outcome-less (22 d stale); 4 are live
TOTAL_VERIFICATIONS_OBSERVED	14 (11 claim verifications + 3 closure verdicts)	claims.db::verifications · closure-ledger.jsonl	1 of the 11 is PARTIAL (evidence hash moved under the claim)
TOTAL_LESSONS_OBSERVED	29 (12 unratified + 13 scars) + 4 scar candidates + 346 RSI atoms	UNRATIFIED-LESSONS-LEDGER.jsonl · scars/ · atoms.jsonl	atoms: 318 verified / 28 rejected; every verification provisional
TOTAL_POLICY_CHANGES_OBSERVED	0	no artifact exists	nothing to bind — the object class is absent

6.2 Which of these is CHRON

Candidate identity	Mark	Evidence, or the absence of it
Scheduler	[x]	39-job registry + delivery ledger + 30 systemd timers + ~45 crontab lines, all live today. But none of them is CHRON's — CHRON owns no trigger. Scheduler capability exists in the substrate; CHRON is a client of it, not the owner.
Relevance Engine	[x]	chron.py ranks by urgency × consequence × actionability × confidence with a non-linear urgency curve and drops expired events; the card gate selects 9 of 18 signals and refuses a bad card. Scoring is real and executable. Its trigger is not.
Claim-State Engine	[]	A real claim-state store exists — claims.db, 21 claims, 11 verifications, append-only triggers, supersedes_claim_id. It is a separate organ ; no CHRON-owned claim-state object was observed, and nothing in CHRON reads it automatically.
Temporal Reconciliation Layer	[]	This is the declared identity. Observed: a gate script that reconciles <i>declared boundaries against artifacts</i> (6 PASS / 3 UNBUILT) and runs only when a human types its name. Declared-vs-observed reconciliation as a daily job does not exist — which is why D1-D8 were still open today.
Learning System	[]	The loop runs (14 cycles, 346 atoms) and does not close: 8 of 14 cycles inhaled without exhaling, latest 4 consecutive cycles INHALE_ONLY, all 7 baselines PENDING, 7

CURRENT CHRON IS: Scheduler ☐ (*substrate, not owned*) · Relevance Engine ☒ · Claim-State Engine ☐ · Temporal Reconciliation Layer ☐ · Learning System ☐. **Two of five, and both of the two are the parts that do not require the loop to close.**

FINAL VERDICT — five items only

1 · CURRENT REALITY

A date store, a gate, and a card lane that has never fired.

What is alive on this machine under the name CHRON: **5 canonically-sourced events** in one store; a **spine gate** that executes clean and self-tests its own ability to reject (6 PASS · 0 FAIL · 3 UNBUILT); a **relevance engine** (`chron.py`) that computes every countdown at render and expires what has passed; and a **two-per-day card lane** with two rendered artifacts and **zero executions**.

What is alive around it and often mistaken for it: a 39-job scheduler (13 enabled), an RSI learning loop firing every 6 hours, a daily attention closure at 23:45, a claim ledger holding 21 claims and 11 verifications, and a briefing lane that produced a real PDF this morning. CHRON wrote the contract for the card; CHRON does not own the scheduler that would run it, and nothing schedules CHRON's own scripts.

And one thing is definitively dead: the world-model prediction stream — 792 rows, 5 distinct actions, no outcome field, last write 22 days ago, no reader.

2 · ACTIVE TASKS

Seven carry a live trigger. Three are armed but never fired. Two are dormant code.

Live and evidenced: `docforge-edition-daily` (06:00 → ran 09:59 today) · `attention-closure` (23:45 → last 09:17) · `aaa-rsi-loop` (every 6 h → 12:07 today) · `aaa-mesh-health` (03:17) · the root `crontab` loops (metabolism 30 min, reality-pulse 15 min, cockpit probe 15 min, well intake 30 min, `arifFlow` digest 22:00) · `systemd` timers including `vault999-backup`, `triadic-snapshot`, `arifos-reality`.

Armed, never executed: the three ALPHA-ZEN jobs — 07:15 (missed by construction: created 10:02), 14:00 (next, in ~2 h), 21:15.

Dormant code with no trigger: `chron.py` and `chron_spine_gate.py` — both run correctly, both invoked only by hand or from inside another job's prompt text.

3 · MISSING LOOPS

Three, and they are the three that would let CHRON learn.

L6 Outcome → Calibration — nothing scores anything. No Brier, log-loss, or ECE artifact exists. The single recorded failure (`REALITY DISAGREED`, `silent-count` predicted 7, observed 6) produced no calibration entry. **L8 Lesson → Policy** — 12 lessons sit unratified; no policy-candidate object exists, so there is nothing a lesson could be promoted into. **L9 Policy → Attention** — attention is generated from infrastructure gaps, never from an adopted policy; a change in what matters does not change what is looked at.

Two further arrow-level gaps sit behind these, both reported UNBUILT by CHRON's own gate: **EXPERIENCE → MEMORY** (a collection exists with no payload schema) and **MEMORY → POLICY**.

4 · HIGHEST DRIFT

Three enabled jobs, a full pipeline, a named room — and zero receipts.

D1, highest. The ALPHA-ZEN contract is complete: gate that refuses, canonical clock, schema, audience rule, memory boundary, delivery target, and two artifacts rendered by hand this morning. All three jobs read `enabled: true` · `state: scheduled` · `last_run_at: null` · `completed: 0`, and all three were created at 10:02 today — **after the 07:15 slot they were built for**. Nothing here is broken; nothing here has run. The gap is a missing receipt, and the window to close it is today at 14:00.

D5, structural. 35,743 execution receipts, **100 % without a trace_id**. The federation already named this defect in doctrine; today it is measurable in the running system. Not a missing log — a missing causal join.

D3, silent. A 792-row prediction stream that looks healthy by every surface check (valid JSONL, growing file, distinct ids) and stopped writing 22 days ago, with no reader and no outcome field. Its shape is the warning: artifacts that look fine while the thing behind them has ceased.

Also open: D8 — the ALPHA-ZEN audit still lists four blockers that were repaired in code this morning with no supersession recorded. The artifact that taught this lesson is now an example of it.

5 · NEXT LOWEST-RISK WITNESS ACTION

Witness the 14:00 fire — capture the receipt, not the intent.

Action: at 14:00 MYT today, read-only, record for the ALPHA-ZEN body-check job: (a) `last_run_at` and `last_status` from `jobs.json`; (b) the delivery row from `deliveries.db`; (c) any artifact `mtime/hash` produced; (d) the Telegram `message_id` if the receipt carries one. Then file the result as **DELIVERED / OBSERVED / ACKNOWLEDGED**, separately — never as one word.

Why this one: zero mutation, zero new architecture, no F13 decision required, and it is the only step that converts the single largest drift from *armed but unproven* into either a receipt or a named fault. It also closes loop L1 with real evidence instead of a scheduled promise — and it directly answers the question Arif was asking this morning about a PDF he could not find: **whether a scheduled lane produced anything at all, or only a schedule.**

Named hazard for that run: the two prior card renders happened by hand. If 14:00 produces a message without a `message_id`, that is the same D5/D6 class already recorded against the digest lane — channel proven, event unbound. Report it that way; do not read a delivery row as proof the room saw it.

Standing constraints observed while writing this

OBSERVE_ONLY: no file outside this artifact's own directory was created or modified. **DECLARATION ≠ MECHANISM:** every capability claim above is bound to a path that was opened or a command that was executed during this pass.

WITNESS BEFORE MUTATION: the eight drifts are recorded, not repaired — repair is a separate, authorised action.

Naming honesty: this is a **Lane B RECEIPT**, not a SEAL. `seal_allowed=false` is reported here from the day's own prior artifact, not re-probed by this seat; no judge hash, no witness, no cryptographic actor binding is claimed. The word SEAL does not appear as a verdict anywhere in this document.

CHRON::REALITY_MAPPER::v1 · Authority: ARIF (F13) · Mode: OBSERVE_ONLY / AUDIT · Render: 2026-09-18 12:11 MYT · Seat: HERMES (i-ARIF), KVM8 · Class: Lane B RECEIPT · Mutations: 0

Probe set executed this pass: `jobs.json` (39 jobs) · `deliveries.db` · `systemctl list-timers` · `crontab -l + /etc/cron.d` · `claims.db` (4 tables) · `closure-ledger.jsonl` · `attention-ledger.jsonl` + `snapshot` · `rsi/state/*` (atoms, loop-ledger, consequence, baselines) · `UNRATIFIED-LESSONS-LEDGER.jsonl` · `scars/` · `chron_events.json` · `cycles.jsonl` · `cards/` · `world-model-lite.jsonl` (792 rows parsed) · `ariffow_receipts.jsonl` (35,743 rows parsed) · `chron_spine_gate.py` **executed** · Qdrant collection list · KVM4 /health + cron API.

DITEMPA BUKAN DIBERI